

Reppeto530Week3

December 17, 2023

1 Chapter 1

Brian Reppeto 530 Prof. Jim Week 3 HW

```
[35]: # Import and download the code from the github repo

from os.path import basename, exists

def download(url):
    filename = basename(url)
    if not exists(filename):
        from urllib.request import urlretrieve

        local, _ = urlretrieve(url, filename)
        print("Downloaded " + local)

download("https://github.com/AllenDowney/ThinkStats2/raw/master/code/
↳thinkstats2.py")
download("https://github.com/AllenDowney/ThinkStats2/raw/master/code/thinkplot.
↳py")

[36]: # This part is to download the specific files needed for the assignment

download("https://github.com/AllenDowney/ThinkStats2/raw/master/code/nsfg.py")

download("https://github.com/AllenDowney/ThinkStats2/raw/master/code/
↳2002FemPreg.dct")
download(
    "https://github.com/AllenDowney/ThinkStats2/raw/master/code/2002FemPreg.dat.
↳gz"
)

[37]: # Importing the .py file to work with in the assignment

import nsfg
```

```
[79]: # Reading the .py file and setting it = to preg
preg = nsfg.ReadFemPreg()
```

1.1 Exercise 1-1

Select the `birthord` column, print the value counts, and compare to results published in the [codebook](#)

```
[80]: # Using the preg DF, select the "birthord" column and count the number of times
      ↪ each occur.
      # Then sort the output based on the cloumn
preg.birthord.value_counts().sort_index()
```

```
[80]: birthord
1.0    4413
2.0    2874
3.0    1234
4.0     421
5.0     126
6.0      50
7.0      20
8.0       7
9.0       2
10.0      1
Name: count, dtype: int64
```

We can also use `isnull` to count the number of nans.

```
[50]: # Using the preg DF, select the "birthord" column and sum the number of times
      ↪ each the column isnull.

preg.birthord.isnull().sum()
```

```
[50]: 4445
```

Select the `prglngth` column, print the value counts, and compare to results published in the [codebook](#)

```
[51]: # Using the preg DF, select the "prglngth" column and count the number of times
      ↪ each occur.
      # Then sort the output based on the cloumn
preg.prglngth.value_counts().sort_index()
```

```
[51]: prglngth
0      15
1       9
2      78
3     151
```

4	412
5	181
6	543
7	175
8	409
9	594
10	137
11	202
12	170
13	446
14	29
15	39
16	44
17	253
18	17
19	34
20	18
21	37
22	147
23	12
24	31
25	15
26	117
27	8
28	38
29	23
30	198
31	29
32	122
33	50
34	60
35	357
36	329
37	457
38	609
39	4744
40	1120
41	591
42	328
43	148
44	46
45	10
46	1
47	1
48	7
50	2

Name: count, dtype: int64

To compute the mean of a column, you can invoke the `mean` method on a Series. For example, here is the mean birthweight in pounds:

```
[52]: preg.totalwgt_lb.mean()
```

```
[52]: 7.265628457623368
```

Create a new column named `totalwgt_kg` that contains birth weight in kilograms. Compute its mean. Remember that when you create a new column, you have to use dictionary syntax, not dot notation.

```
[57]: # Create a column named "totalwgt_kg" there are 2.2 Kg in a pound so convert
      ↪ the pound column to get the Kg.
      # Then set the new column = to the calc
      # print the mean of the new kg column

preg['totalwgt_kg'] = preg.totalwgt_lb / 2.2
# print (preg.head())
preg.totalwgt_kg.mean()
```

```
[57]: 3.302558389828803
```

`nsfg.py` also provides `ReadFemResp`, which reads the female respondents file and returns a `DataFrame`:

```
[58]: download("https://github.com/AllenDowney/ThinkStats2/raw/master/code/
      ↪ 2002FemResp.dct")
      download("https://github.com/AllenDowney/ThinkStats2/raw/master/code/
      ↪ 2002FemResp.dat.gz")
```

Downloaded 2002FemResp.dct

Downloaded 2002FemResp.dat.gz

```
[59]: resp = nsfg.ReadFemResp()
```

`DataFrame` provides a method `head` that displays the first five rows:

```
[60]: resp.head()
```

```
[60]:   caseid  rscrinf  rdormres  rostscrn  rscreenhisp  rscreenrace  age_a  \
0    2298        1         5         5           1          5.0    27
1    5012        1         5         1           5          5.0    42
2   11586        1         5         1           5          5.0    43
3    6794        5         5         4           1          5.0    15
4     616        1         5         4           1          5.0    20

   age_r  cmbirth  agescrn  ...  pubassis_i  basewgt  adj_mod_basewgt  \
0     27     902      27  ...           0  3247.916977    5123.759559
1     42     718      42  ...           0  2335.279149    2846.799490
2     43     708      43  ...           0  2335.279149    2846.799490
```

3	15	1042	15	...	0	3783.152221	5071.464231
4	20	991	20	...	0	5341.329968	6437.335772

	finalwgt	secu_r	sest	cmintvw	cmlstyr	screentime	intvlngth
0	5556.717241	2	18	1234	1222	18:26:36	110.492667
1	4744.191350	2	18	1233	1221	16:30:59	64.294000
2	4744.191350	2	18	1234	1222	18:19:09	75.149167
3	5923.977368	2	18	1234	1222	15:54:43	28.642833
4	7229.128072	2	18	1233	1221	14:19:44	69.502667

[5 rows x 3087 columns]

Select the `age_r` column from `resp` and print the value counts. How old are the youngest and oldest respondents?

```
[62]: # Using the resp DF, select the "age_r" column and count the number of times
      ↪ each occur.
      # Then sort the output based on the cloumn

      resp.age_r.value_counts().sort_index()
```

```
[62]: age_r
      15    217
      16    223
      17    234
      18    235
      19    241
      20    258
      21    267
      22    287
      23    282
      24    269
      25    267
      26    260
      27    255
      28    252
      29    262
      30    292
      31    278
      32    273
      33    257
      34    255
      35    262
      36    266
      37    271
      38    256
      39    215
      40    256
```

```

41    250
42    215
43    253
44    235
Name: count, dtype: int64

```

We can use the `caseid` to match up rows from `resp` and `preg`. For example, we can select the row from `resp` for `caseid` 2298 like this:

```
[63]: resp[resp.caseid==2298]
```

```

[63]:   caseid  rscrinf  rdormres  rostscrn  rscreenhisp  rscreenrace  age_a  \
0    2298         1         5         5         1         5.0    27

      age_r  cmbirth  agescrn  ...  pubassis_i      basewgt  adj_mod_basewgt  \
0     27      902      27  ...         0  3247.916977      5123.759559

      finalwgt  secu_r  sest  cmintvw  cmlstyr  screentime  intvlngth
0  5556.717241         2   18    1234    1222    18:26:36   110.492667

[1 rows x 3087 columns]

```

And we can get the corresponding rows from `preg` like this:

```
[64]: preg[preg.caseid==2298]
```

```

[64]:   caseid  pregordr  howpreg_n  howpreg_p  moscurrp  nowprgdk  pregend1  \
2610    2298         1        NaN        NaN        NaN        NaN        6.0
2611    2298         2        NaN        NaN        NaN        NaN        6.0
2612    2298         3        NaN        NaN        NaN        NaN        6.0
2613    2298         4        NaN        NaN        NaN        NaN        6.0

      pregend2  nbrnaliv  multbrth  ...  religion_i  metro_i      basewgt  \
2610        NaN        1.0        NaN  ...         0         0  3247.916977
2611        NaN        1.0        NaN  ...         0         0  3247.916977
2612        NaN        1.0        NaN  ...         0         0  3247.916977
2613        NaN        1.0        NaN  ...         0         0  3247.916977

      adj_mod_basewgt      finalwgt  secu_p  sest  cmintvw  totalwgt_lb  \
2610      5123.759559  5556.717241         2   18        NaN        6.8750
2611      5123.759559  5556.717241         2   18        NaN        5.5000
2612      5123.759559  5556.717241         2   18        NaN        4.1875
2613      5123.759559  5556.717241         2   18        NaN        6.8750

      totalwgt_kg
2610      3.125000
2611      2.500000
2612      1.903409

```

```
2613      3.125000
```

```
[4 rows x 245 columns]
```

How old is the respondent with caseid 1?

```
[67]: # Use the DF resp and then select the DF resp with column caseid that =1 then
      ↪return the column age_r to return the age of
      # the pregnancy with caseid =1

      resp[resp.caseid==1].age_r
```

```
[67]: 1069      44
      Name: age_r, dtype: int64
```

What are the pregnancy lengths for the respondent with caseid 2298?

```
[77]: # Use the DF preg and then select the DF preg with column caseid that = 2298
      ↪then return the column prglngth to return the age of
      # the pregnancy with caseid = 2298

      preg[preg.caseid==2298].prglngth
```

```
[77]: 2610      40
      2611      36
      2612      30
      2613      40
      Name: prglngth, dtype: int64
```

What was the birthweight of the first baby born to the respondent with caseid 5012?

```
[78]: # Use the DF preg and then select the DF preg with column caseid that = 5012
      ↪then return the column birthwgt_lb
      # to return the birth weight. This respondent only had one pregnancy.

      preg[preg.caseid==5012].birthwgt_lb
```

```
[78]: 5515      6.0
      Name: birthwgt_lb, dtype: float64
```

1.2 Exercise 1-2

```
[81]: # Reading the .py file and setting it = to resp

      resp=nsfg.ReadFemResp()
```

```
[84]: # Print the value count for this variable
```

```
resp.pregnum.value_counts()
```

```
[84]: pregnum
```

```
0      2610
2      1432
1      1267
3      1110
4       611
5       305
6       150
7        80
8        40
9        21
10        9
11         3
12         2
14         2
19         1
```

```
Name: count, dtype: int64
```

```
[87]: # Print the value count to compare to the resp DF. The counts do not align.
```

```
preg.pregnum.value_counts()
```

```
[87]: pregnum
```

```
3      3330
2      2864
4      2444
5      1525
1      1267
6       900
7       560
8       320
9       189
10       90
11       33
14       28
12       24
19       19
```

```
Name: count, dtype: int64
```

```
[100]: # Test the cross-validation via caseid.
```

```
# Print the resp df caseid ==2 identifying the caseid and the pregnum fields to
↳compare cross-val.
```



```

resp_rows = resp[resp.caseid == 2]
print(resp_rows[['caseid', 'pregnum']])

# Print the preg df caseid ==2 identifying the caseid and the pregnum fields to
↳ compare cross-val
preg_rows = preg[preg.caseid == 2]
print(preg_rows[['caseid', 'pregnum']])

# both df cross-valid for pregnancy number

```

	caseid	pregnum
7010	2	3
	caseid	pregnum
2	2	3
3	2	3
4	2	3

[104]: *# This maps the caseid to index*

```
nsfg.MakePregMap(resp)
```

```

[104]: defaultdict(list,
    {2298: [0],
     5012: [1],
     11586: [2],
     6794: [3],
     616: [4],
     845: [5],
     10333: [6],
     855: [7],
     8656: [8],
     3566: [9],
     5917: [10],
     9200: [11],
     6320: [12],
     11700: [13],
     7354: [14],
     3697: [15],
     4881: [16],
     5862: [17],
     8542: [18],
     2054: [19],
     3719: [20],
     11740: [21],
     11343: [22],
     7075: [23],

```

5422: [24],
2178: [25],
8358: [26],
5083: [27],
1545: [28],
5656: [29],
9334: [30],
5507: [31],
611: [32],
4260: [33],
11767: [34],
5573: [35],
11901: [36],
8975: [37],
5267: [38],
910: [39],
4463: [40],
8954: [41],
1814: [42],
7011: [43],
4057: [44],
7081: [45],
5499: [46],
6551: [47],
9242: [48],
11408: [49],
7168: [50],
2339: [51],
4138: [52],
8785: [53],
1511: [54],
2240: [55],
8190: [56],
5278: [57],
469: [58],
5351: [59],
10633: [60],
11058: [61],
3329: [62],
3814: [63],
4966: [64],
4327: [65],
6908: [66],
9302: [67],
4627: [68],
8085: [69],
9555: [70],

1821: [71],
319: [72],
10295: [73],
8269: [74],
4693: [75],
6370: [76],
8628: [77],
2171: [78],
581: [79],
8065: [80],
8081: [81],
2286: [82],
7758: [83],
2401: [84],
4795: [85],
6145: [86],
10463: [87],
547: [88],
2522: [89],
997: [90],
1786: [91],
4308: [92],
3800: [93],
7599: [94],
9866: [95],
5338: [96],
10289: [97],
2281: [98],
4150: [99],
3679: [100],
1810: [101],
7369: [102],
5243: [103],
1014: [104],
5455: [105],
8417: [106],
5296: [107],
7483: [108],
6397: [109],
7430: [110],
545: [111],
8661: [112],
3700: [113],
12031: [114],
5288: [115],
9846: [116],
2137: [117],

11470: [118],
3349: [119],
12014: [120],
1703: [121],
3391: [122],
6953: [123],
8270: [124],
175: [125],
9680: [126],
7598: [127],
6716: [128],
251: [129],
12109: [130],
7497: [131],
12219: [132],
1877: [133],
8366: [134],
785: [135],
5321: [136],
8062: [137],
3821: [138],
2749: [139],
6963: [140],
7498: [141],
435: [142],
10069: [143],
9849: [144],
3372: [145],
9527: [146],
5630: [147],
5078: [148],
9810: [149],
11421: [150],
9523: [151],
6281: [152],
5365: [153],
12434: [154],
7114: [155],
6441: [156],
12044: [157],
6032: [158],
2885: [159],
11333: [160],
5343: [161],
1798: [162],
2788: [163],
9726: [164],

5652: [165],
3508: [166],
2460: [167],
505: [168],
8617: [169],
11603: [170],
5479: [171],
2963: [172],
8389: [173],
10242: [174],
2336: [175],
11484: [176],
5539: [177],
8181: [178],
10860: [179],
8195: [180],
6832: [181],
2678: [182],
4790: [183],
3570: [184],
8817: [185],
9927: [186],
10844: [187],
11185: [188],
12267: [189],
9133: [190],
9516: [191],
1944: [192],
2587: [193],
6896: [194],
12205: [195],
7069: [196],
2257: [197],
2456: [198],
4500: [199],
5270: [200],
118: [201],
2077: [202],
7156: [203],
11823: [204],
10367: [205],
9227: [206],
5644: [207],
4200: [208],
12286: [209],
6151: [210],
5068: [211],

61: [212],
116: [213],
2196: [214],
8004: [215],
12400: [216],
1182: [217],
5666: [218],
10016: [219],
2218: [220],
9587: [221],
2200: [222],
4527: [223],
7770: [224],
11777: [225],
8036: [226],
1349: [227],
9983: [228],
2808: [229],
8100: [230],
3026: [231],
10803: [232],
781: [233],
2899: [234],
1561: [235],
5547: [236],
112: [237],
5378: [238],
9974: [239],
1910: [240],
1089: [241],
8196: [242],
1531: [243],
37: [244],
1045: [245],
3598: [246],
9658: [247],
5924: [248],
1518: [249],
11030: [250],
9376: [251],
7543: [252],
6178: [253],
12428: [254],
10113: [255],
5720: [256],
7053: [257],
5166: [258],

10261: [259],
12494: [260],
3660: [261],
1774: [262],
5731: [263],
8413: [264],
10201: [265],
12482: [266],
12215: [267],
11371: [268],
5176: [269],
8544: [270],
4921: [271],
11874: [272],
7799: [273],
5774: [274],
2838: [275],
9190: [276],
5540: [277],
2399: [278],
671: [279],
3319: [280],
2027: [281],
901: [282],
6295: [283],
9969: [284],
7848: [285],
12382: [286],
4233: [287],
1424: [288],
12550: [289],
1612: [290],
3953: [291],
765: [292],
1100: [293],
9667: [294],
4080: [295],
5915: [296],
9254: [297],
6718: [298],
9955: [299],
197: [300],
5323: [301],
9179: [302],
227: [303],
10456: [304],
7965: [305],

11919: [306],
3478: [307],
7926: [308],
6428: [309],
5480: [310],
10535: [311],
12504: [312],
9972: [313],
11747: [314],
1611: [315],
5688: [316],
4858: [317],
6140: [318],
11834: [319],
7628: [320],
4174: [321],
8674: [322],
10140: [323],
3402: [324],
385: [325],
1440: [326],
4832: [327],
8079: [328],
12422: [329],
8561: [330],
4999: [331],
6332: [332],
5478: [333],
10885: [334],
7837: [335],
10029: [336],
4006: [337],
7984: [338],
657: [339],
11958: [340],
1932: [341],
11613: [342],
11212: [343],
1017: [344],
1481: [345],
7046: [346],
6434: [347],
9859: [348],
943: [349],
3728: [350],
9207: [351],
8050: [352],

3447: [353],
11785: [354],
7008: [355],
2297: [356],
1654: [357],
8737: [358],
6594: [359],
2951: [360],
4153: [361],
6456: [362],
1752: [363],
12107: [364],
1707: [365],
6715: [366],
7883: [367],
954: [368],
7569: [369],
4292: [370],
10652: [371],
10247: [372],
12495: [373],
6329: [374],
10105: [375],
4372: [376],
7513: [377],
6782: [378],
7627: [379],
10205: [380],
12193: [381],
8080: [382],
11132: [383],
4885: [384],
10364: [385],
7567: [386],
11774: [387],
11152: [388],
12389: [389],
1315: [390],
7732: [391],
5898: [392],
5022: [393],
11382: [394],
8082: [395],
11858: [396],
11681: [397],
12113: [398],
458: [399],

3934: [400],
5234: [401],
9829: [402],
8743: [403],
10058: [404],
1833: [405],
8465: [406],
9518: [407],
8918: [408],
12344: [409],
393: [410],
2798: [411],
4266: [412],
1147: [413],
541: [414],
11898: [415],
8148: [416],
3463: [417],
5306: [418],
418: [419],
8675: [420],
12448: [421],
5850: [422],
8252: [423],
5842: [424],
7915: [425],
10189: [426],
4069: [427],
1853: [428],
258: [429],
8840: [430],
9842: [431],
7694: [432],
9861: [433],
4258: [434],
1418: [435],
6236: [436],
585: [437],
8334: [438],
6205: [439],
4177: [440],
2873: [441],
10694: [442],
3457: [443],
1499: [444],
5295: [445],
163: [446],

3835: [447],
5160: [448],
12468: [449],
10084: [450],
2306: [451],
6074: [452],
3039: [453],
5107: [454],
7931: [455],
399: [456],
1122: [457],
6546: [458],
2915: [459],
7066: [460],
2213: [461],
2059: [462],
8602: [463],
11975: [464],
3933: [465],
8966: [466],
6859: [467],
12162: [468],
8909: [469],
8613: [470],
10309: [471],
10874: [472],
217: [473],
12072: [474],
11936: [475],
7767: [476],
4923: [477],
11229: [478],
8819: [479],
5405: [480],
8117: [481],
2404: [482],
1963: [483],
5350: [484],
11611: [485],
9449: [486],
10945: [487],
5514: [488],
77: [489],
3811: [490],
22: [491],
10564: [492],
1305: [493],

2415: [494],
4946: [495],
11814: [496],
11445: [497],
5205: [498],
8478: [499],
3054: [500],
10595: [501],
10717: [502],
9696: [503],
7535: [504],
3682: [505],
6791: [506],
1491: [507],
135: [508],
218: [509],
5126: [510],
9283: [511],
10144: [512],
11985: [513],
9963: [514],
9585: [515],
5025: [516],
11710: [517],
832: [518],
8538: [519],
8802: [520],
2104: [521],
8138: [522],
522: [523],
11399: [524],
2959: [525],
374: [526],
6753: [527],
9506: [528],
2176: [529],
6897: [530],
10259: [531],
12190: [532],
11672: [533],
1228: [534],
4606: [535],
8458: [536],
12067: [537],
5048: [538],
2700: [539],
8307: [540],

1734: [541],
10561: [542],
10335: [543],
1422: [544],
1532: [545],
7085: [546],
3269: [547],
3251: [548],
3904: [549],
744: [550],
6640: [551],
8887: [552],
7059: [553],
4096: [554],
12276: [555],
140: [556],
8077: [557],
2406: [558],
5485: [559],
11839: [560],
5030: [561],
10492: [562],
8054: [563],
11144: [564],
4855: [565],
5329: [566],
3966: [567],
9192: [568],
2673: [569],
9675: [570],
941: [571],
6159: [572],
2795: [573],
745: [574],
4346: [575],
219: [576],
1629: [577],
9476: [578],
1597: [579],
95: [580],
8920: [581],
410: [582],
8671: [583],
4289: [584],
7741: [585],
6393: [586],
5403: [587],

751: [588],
4743: [589],
6093: [590],
2983: [591],
12156: [592],
11450: [593],
2605: [594],
2667: [595],
7343: [596],
4889: [597],
11935: [598],
6699: [599],
7509: [600],
497: [601],
11029: [602],
12089: [603],
3281: [604],
9113: [605],
10025: [606],
9348: [607],
12023: [608],
3865: [609],
5665: [610],
2422: [611],
3458: [612],
12507: [613],
10043: [614],
11161: [615],
3545: [616],
3146: [617],
2394: [618],
5996: [619],
3541: [620],
3979: [621],
5327: [622],
4810: [623],
859: [624],
12568: [625],
8139: [626],
9282: [627],
8941: [628],
4241: [629],
5801: [630],
10723: [631],
9394: [632],
546: [633],
7294: [634],

1082: [635],
11285: [636],
9515: [637],
8184: [638],
7698: [639],
4744: [640],
11896: [641],
3094: [642],
10664: [643],
2164: [644],
8242: [645],
1266: [646],
5523: [647],
7969: [648],
1773: [649],
11197: [650],
10777: [651],
6810: [652],
9271: [653],
9223: [654],
7715: [655],
12378: [656],
5538: [657],
8193: [658],
5951: [659],
4093: [660],
2521: [661],
3907: [662],
4565: [663],
3828: [664],
1892: [665],
7882: [666],
10560: [667],
10067: [668],
991: [669],
3738: [670],
10812: [671],
6027: [672],
4231: [673],
1894: [674],
9429: [675],
6694: [676],
11578: [677],
2128: [678],
8066: [679],
627: [680],
2535: [681],

1169: [682],
1171: [683],
734: [684],
8276: [685],
632: [686],
9624: [687],
9177: [688],
2439: [689],
12180: [690],
12564: [691],
7333: [692],
12294: [693],
738: [694],
9355: [695],
11471: [696],
723: [697],
11941: [698],
12493: [699],
544: [700],
10708: [701],
12074: [702],
4562: [703],
4160: [704],
8350: [705],
173: [706],
9797: [707],
6339: [708],
8532: [709],
9062: [710],
9953: [711],
11994: [712],
12556: [713],
225: [714],
7263: [715],
72: [716],
7703: [717],
365: [718],
8906: [719],
8089: [720],
4154: [721],
10103: [722],
5268: [723],
7861: [724],
1366: [725],
535: [726],
7331: [727],
8715: [728],

8279: [729],
12284: [730],
9: [731],
786: [732],
2280: [733],
12008: [734],
9714: [735],
6398: [736],
6083: [737],
3742: [738],
12356: [739],
7588: [740],
11215: [741],
10006: [742],
10152: [743],
7633: [744],
9526: [745],
6705: [746],
8702: [747],
3678: [748],
6840: [749],
7083: [750],
7507: [751],
4184: [752],
1556: [753],
6889: [754],
4540: [755],
4619: [756],
8325: [757],
2076: [758],
11665: [759],
1465: [760],
5144: [761],
9887: [762],
6196: [763],
12018: [764],
9059: [765],
39: [766],
6417: [767],
9882: [768],
5747: [769],
5948: [770],
10974: [771],
3555: [772],
10836: [773],
2609: [774],
11296: [775],

10530: [776],
9350: [777],
1313: [778],
11897: [779],
11857: [780],
890: [781],
5636: [782],
5597: [783],
9478: [784],
11601: [785],
3115: [786],
10193: [787],
12181: [788],
7992: [789],
5501: [790],
9612: [791],
2601: [792],
4631: [793],
12238: [794],
144: [795],
7544: [796],
2479: [797],
5167: [798],
10041: [799],
12451: [800],
7889: [801],
3941: [802],
12324: [803],
4123: [804],
11413: [805],
2825: [806],
238: [807],
548: [808],
11021: [809],
10083: [810],
6299: [811],
3523: [812],
6977: [813],
9139: [814],
7643: [815],
6051: [816],
7153: [817],
907: [818],
10454: [819],
237: [820],
9063: [821],
5607: [822],

3283: [823],
7921: [824],
6348: [825],
6303: [826],
7142: [827],
3862: [828],
11092: [829],
2048: [830],
1476: [831],
2239: [832],
2900: [833],
2629: [834],
9613: [835],
10439: [836],
2474: [837],
1768: [838],
7219: [839],
5712: [840],
8442: [841],
463: [842],
4558: [843],
3871: [844],
11532: [845],
21: [846],
8240: [847],
3743: [848],
1560: [849],
7254: [850],
2641: [851],
6916: [852],
4086: [853],
7695: [854],
11753: [855],
2842: [856],
2818: [857],
7802: [858],
918: [859],
717: [860],
2807: [861],
11910: [862],
9998: [863],
945: [864],
1701: [865],
11953: [866],
4238: [867],
2130: [868],
1099: [869],

2593: [870],
6125: [871],
9255: [872],
5064: [873],
2534: [874],
3188: [875],
1829: [876],
8239: [877],
4986: [878],
4820: [879],
10528: [880],
1939: [881],
6489: [882],
7854: [883],
2596: [884],
1138: [885],
10404: [886],
3656: [887],
12252: [888],
8756: [889],
1645: [890],
6856: [891],
8907: [892],
10918: [893],
4317: [894],
7386: [895],
12114: [896],
6161: [897],
44: [898],
10795: [899],
1995: [900],
10683: [901],
1790: [902],
10690: [903],
3178: [904],
10534: [905],
8923: [906],
10079: [907],
7952: [908],
10906: [909],
3140: [910],
5832: [911],
11068: [912],
4303: [913],
555: [914],
7123: [915],
867: [916],

11102: [917],
1699: [918],
6553: [919],
1705: [920],
10150: [921],
3116: [922],
11372: [923],
7804: [924],
2966: [925],
10380: [926],
8639: [927],
10729: [928],
6637: [929],
3159: [930],
3218: [931],
4027: [932],
9466: [933],
132: [934],
1153: [935],
9634: [936],
6964: [937],
6268: [938],
710: [939],
3003: [940],
12010: [941],
8934: [942],
10926: [943],
8865: [944],
11776: [945],
1575: [946],
2033: [947],
6266: [948],
10168: [949],
1906: [950],
7504: [951],
494: [952],
6403: [953],
3111: [954],
7001: [955],
9668: [956],
9900: [957],
10821: [958],
1581: [959],
8255: [960],
1570: [961],
8783: [962],
4628: [963],

```

3249: [964],
7756: [965],
1062: [966],
10072: [967],
8075: [968],
5513: [969],
11105: [970],
11078: [971],
8650: [972],
7934: [973],
11032: [974],
12240: [975],
952: [976],
4893: [977],
897: [978],
3597: [979],
2155: [980],
3770: [981],
3529: [982],
5535: [983],
6764: [984],
6967: [985],
5703: [986],
9447: [987],
12226: [988],
2165: [989],
3486: [990],
10160: [991],
11988: [992],
3803: [993],
2985: [994],
2005: [995],
8841: [996],
11377: [997],
4910: [998],
4689: [999],
...})

```

1.3 Exercise 2-1

Evening News

If I were to summarize on the evening news whether first babies arrive late, I would use charts and graphs as they are easier for people to get the overall meaning of the data. Specifically, I would use a Histogram that shows the pregnancy in weeks with the first child being highlighted compared to all other births. I would also use summarizing distributions to show the details between the first and other births.

1.4 Anxious Patient

Reassuring an anxious patient would require using stats that emphasize differences in the situations or scenario to put the differences into context. A few examples of these would be Effect Size to understand the difference in the scenarios, Odds Ratios to compare the odds between the scenarios, also I think Summarizing Distributions would help explain variability and probabilities between scenarios that patients would be able to understand in layman's terms.

1.5 Straight Dope question

Based on the statistical measurements in chapter 2, it is apparent that the belief in first babies arriving late is not statistically supported. The histogram (figure 2-5) depicting pregnancy lengths between first born and all others does not support this hypothesis and actually shows a mean where other born were at a much higher clip than that of first born. The summarization of distributions further underscores the subtle variations between first-time mothers and those with subsequent pregnancies, suggesting that the expectation of first babies arriving late may not hold true across the board. While individual experiences may vary, the comprehensive analysis of pregnancy data does not decisively support the widespread notion that first babies are predisposed to tardy arrivals. The intricacies of childbirth, influenced by numerous factors, warrant a nuanced understanding that goes beyond generalized assumptions.

1.6 Exercise 2-4

```
[108]: # Create a DF for the live births from preg
```

```
live=preg[preg.outcome==1]
```

```
# create a division of the live df to subset first born from other born
```

```
first_born=live[live.birthord==1]
```

```
other_born=live[live.birthord!=1]
```

```
[116]: # subselect each fb mean and other mean then print the mean of each in lb's to
      ↪ determine if first babies are
      # lighter or heavier. The calc shows FB as lighter than other born in terms of
      ↪ the mean.
```

```
fb_mean=first_born.totalwgt_lb.mean()
```

```
other_mean=other_born.totalwgt_lb.mean()
```

```
print(f"first born: {fb_mean} mean lb's")
```

```
print(f"other born: {other_mean} mean lb's")
```

```
first born: 7.201094430437772 mean lb's
```

```
other born: 7.325855614973262 mean lb's
```

```
[117]: # define Cohen's d effect
```

```
#import library
```

```

import numpy as np

def CohenEffectSize(group1, group2):
    """Computes Cohen's effect size for two groups.

    group1: Series or DataFrame
    group2: Series or DataFrame

    returns: float if the arguments are Series;
             Series if the arguments are DataFrames
    """
    diff = group1.mean() - group2.mean()

    var1 = group1.var()
    var2 = group2.var()
    n1, n2 = len(group1), len(group2)

    pooled_var = (n1 * var1 + n2 * var2) / (n1 + n2)
    d = diff / np.sqrt(pooled_var)
    return d

```

```

[118]: # Calculate Cohen's d to compare the first born birth weight in lb's to other_
        ↪ born birth weights in lb's.

```

```

CohenEffectSize(first_born.totalwgt_lb, other_born.totalwgt_lb)

```

```

[118]: -0.088672927072602

```

```

[120]: # subselect each fb length mean and other mean then print the mean of each in_
        ↪ length to determine if first babies are
        # longer or shorter. The calc shows FB as shorter than other born in terms of_
        ↪ the mean.

```

```

fb_ln_mean=first_born.prglngh.mean()
other_ln_mean=other_born.prglngh.mean()
print(f"first born length: {fb_mean} mean")
print(f"other born length: {other_mean} mean")

```

```

first born length: 7.201094430437772 mean
other born length: 7.325855614973262 mean

```

```

[119]: # Calculate Cohen's d to compare the first born birth length to other born_
        ↪ birth length.

```

```

CohenEffectSize(first_born.prglngh, other_born.prglngh)

```

```

[119]: 0.02887904465444988

```


1.7 Response

The primary difference between the two Cohen's d values is their magnitude and the direction of the effect. The first value (-0.08867) (Birth weight) suggests a slightly larger effect size in the negative direction, while the second value (0.028879) (Birth length) suggests a smaller effect size in the positive direction. The negative Birth weight effect suggests that first born babies are lighter than other born the low negative effect suggests only a small difference in the mean of first born, The positive, yet smaller effect suggests there is less of a mean difference in the first born length compared to the other born length.

[]: