

# Brian\_Reppeto\_DSC550\_Week\_11

May 26, 2024

## 0.0.1 DSC 550 Week :

### Activity 11.2

Author: Brian Reppeto 5/22/2024


```
[12]: import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense
from tensorflow.keras.utils import to_categorical
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
```

```
[13]: # load the MNIST dataset

(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

```
[14]: # display the first five images

for i in range(5):
    plt.subplot(1, 5, i+1)
    plt.imshow(X_train[i], cmap='gray')
    plt.title(y_train[i])
    plt.axis('off')
plt.show()
```



```
[15]: # process the data
```

```
X_train = X_train.reshape(X_train.shape[0], 28, 28, 1).astype('float32') / 255
X_test = X_test.reshape(X_test.shape[0], 28, 28, 1).astype('float32') / 255
```

```
[16]: # one-hot encode the labels
```

```
y_train = to_categorical(y_train)
y_test = to_categorical(y_test)
```

```
[17]: # build the CNN model
```

```
model = Sequential([
    Conv2D(32, kernel_size=(3, 3), activation='relu', input_shape=(28, 28, 1)),
    MaxPooling2D(pool_size=(2, 2)),
    Conv2D(64, kernel_size=(3, 3), activation='relu'),
    MaxPooling2D(pool_size=(2, 2)),
    Flatten(),
    Dense(128, activation='relu'),
    Dense(10, activation='softmax')
])
```

```
[18]: # compile the model
```

```
model.compile(optimizer='adam', loss='categorical_crossentropy',
    ↪metrics=['accuracy'])
```

```
[19]: # train the model
```

```
model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=10,
    ↪batch_size=200, verbose=2)
```

Epoch 1/10

300/300 - 6s - loss: 0.2475 - accuracy: 0.9294 - val\_loss: 0.0602 -  
val\_accuracy: 0.9823 - 6s/epoch - 19ms/step

Epoch 2/10

300/300 - 6s - loss: 0.0623 - accuracy: 0.9812 - val\_loss: 0.0450 -  
val\_accuracy: 0.9873 - 6s/epoch - 19ms/step

Epoch 3/10

300/300 - 6s - loss: 0.0441 - accuracy: 0.9866 - val\_loss: 0.0341 -  
val\_accuracy: 0.9889 - 6s/epoch - 19ms/step

Epoch 4/10

300/300 - 6s - loss: 0.0351 - accuracy: 0.9896 - val\_loss: 0.0386 -  
val\_accuracy: 0.9874 - 6s/epoch - 19ms/step

Epoch 5/10

300/300 - 6s - loss: 0.0282 - accuracy: 0.9909 - val\_loss: 0.0296 -  
val\_accuracy: 0.9903 - 6s/epoch - 19ms/step

Epoch 6/10

300/300 - 6s - loss: 0.0218 - accuracy: 0.9934 - val\_loss: 0.0266 -  
val\_accuracy: 0.9918 - 6s/epoch - 19ms/step

```
Epoch 7/10
300/300 - 6s - loss: 0.0188 - accuracy: 0.9942 - val_loss: 0.0287 -
val_accuracy: 0.9907 - 6s/epoch - 19ms/step
Epoch 8/10
300/300 - 6s - loss: 0.0148 - accuracy: 0.9953 - val_loss: 0.0336 -
val_accuracy: 0.9897 - 6s/epoch - 19ms/step
Epoch 9/10
300/300 - 6s - loss: 0.0123 - accuracy: 0.9959 - val_loss: 0.0266 -
val_accuracy: 0.9922 - 6s/epoch - 19ms/step
Epoch 10/10
300/300 - 6s - loss: 0.0104 - accuracy: 0.9966 - val_loss: 0.0329 -
val_accuracy: 0.9896 - 6s/epoch - 19ms/step
```

```
[19]: <keras.callbacks.History at 0x313e6dc90>
```

```
[20]: # evaluate the model

test_loss, test_accuracy = model.evaluate(X_test, y_test, verbose=0)
print(f'Test accuracy: {test_accuracy:.4f}')
```

```
Test accuracy: 0.9896
```


```
[21]: # predict the test set

y_pred = model.predict(X_test)
y_pred_classes = np.argmax(y_pred, axis=1)
y_true = np.argmax(y_test, axis=1)
```

```
313/313 [=====] - 1s 3ms/step
```

```
[22]: # confusion matrix

cm = confusion_matrix(y_true, y_pred_classes)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=range(10))
disp.plot(cmap=plt.cm.Blues)
plt.show()
```



The CNN image classifier trained on the MNIST dataset achieved a test accuracy of approximately 99%. The high accuracy indicates that the model is effective at recognizing handwritten digits. The confusion matrix further illustrates the model's performance, showing a high number of correct predictions across all digit classes. There were a few misclassifications, which demonstrates the model's robustness in distinguishing between different digits.

Overall, the model performs exceptionally well, making it a reliable tool for digit recognition tasks.

[ ]:

[ ]: